

A REPRODUCIBLE SAME-MACHINE COMPARISON OF GRAPH-LAPLACIAN SOLVERS, WITH TWO ROBUSTNESS REFINEMENTS TO LEAN ALGEBRAIC MULTIGRID*

OREN E. LIVNE†

Abstract. Graph-Laplacian systems $L\phi = b$ where L is sparse with m nonzeros underlie spectral clustering, centrality, semi-supervised learning, finite-element analysis, and interior-point network-flow solvers. We present a *reproducible, same-machine study* of graph-Laplacian solvers centered on **LAMG+**—an open, parameter-free Julia re-derivation of Lean Algebraic Multigrid (LAMG) with two targeted robustness refinements—benchmarked under identical conditions against both approximate-Cholesky variants, the degree-aware aggregation of Napov and Notay, and four further state-of-the-art solvers. We establish three facts. (1) *Fair, same-machine benchmarking* against both approximate-Cholesky (AC) variants—the modern near-linear champion—and four other state-of-the-art solvers (hypr/BoomerAMG, PETSc GAMG, pyAMG, CMG). LAMG+ and AC are *complementary peers*: AC is faster on social and citation graphs, LAMG+ on finite-element/structural matrices (where it is the fastest robust solver and the most memory-frugal) and $2.2\times$ faster than the robust AC variant on large graphs. LAMG+ and AC are the only solvers that converge across all 13 test classes, while BoomerAMG, PETSc, pyAMG, and CMG each slow by an order of magnitude or fail to converge off their home turf. (2) *Linear scaling*: LAMG+ is empirically $\mathcal{O}(m)$ over the full 1,711-graph SuiteSparse set (100% converged, median 4 cycles, log–log slope 1.01), verified to 2.4×10^8 nonzeros. (3) *Robustness*: the original MATLAB LAMG was reported non-convergent across the AC benchmark’s test families, yet the *unmodified* LAMG 2.2.1, run on those families with the benchmark’s own solver call, converges on every one—as does LAMG+, to their 10^{-8} tolerance—placing the reported failure in the evaluation, not the algorithm. The two refinements, a strength-of-connection aggregation veto and a selective caliber-2 interpolation, resolve LAMG’s grid-aligned-anisotropy convergence failure (asymptotic convergence factor $\approx 0.99 \rightarrow 0.11$). Both trigger locally and do not substantially increase the complexity. Code and benchmark scripts: <https://github.com/orenlivne/lamgplus>.

Key words. graph Laplacian, algebraic multigrid, aggregation, piecewise-constant interpolation, linear-time solver, approximate Cholesky, adaptive multigrid

AMS subject classifications. 65F08, 65F10, 65N55, 05C50, 90C35

1. Introduction. Let $G = (V, E, w)$ be a connected weighted undirected graph with $|V| = n$ nodes, $|E| = m$ edges, and positive weights $w : E \rightarrow \mathbb{R}^+$. Its Laplacian $L = D - W$ (W the weighted adjacency, D the weighted-degree diagonal) is symmetric positive semidefinite with the constant vector $\mathbf{1}$ spanning its null space; equivalently $\phi^\top L\phi = \sum_{(u,v) \in E} w_{uv}(\phi_u - \phi_v)^2$. We solve the compatible system

$$(1.1) \quad L\phi = b, \quad \mathbf{1}^\top \phi = 0, \quad \mathbf{1}^\top b = 0.$$

Typically $m \ll n^2$ and L is sparse. Our goal is an iterative solver for (1.1) that uses $\mathcal{O}(m)$ storage and $\mathcal{O}(m \log(1/\varepsilon))$ operations for an ε -accurate solution, with *bounded hidden constants on the graphs that arise in applications*—hundreds, not millions. As in the original LAMG work [1], we target good *empirical* wall-clock time performance over a diverse test set rather than provable worst-case complexity on adversarial graphs [2].

1.1. Applications. System (1.1) is fundamental to many computations [2, §2]: elliptic PDEs discretized by finite elements on unstructured grids [3]; interior-point methods for network-flow linear programming, whose inner solve is a weighted Laplacian [4]; electrical flow through a resistor network; spectral clustering, graph embedding, and ancestry discovery in machine learning [5]; and the Fiedler eigenvector, which measures algebraic connectivity

*Received June 19, 2026. Accepted Published online on Recommended by

†Pine Birch Analytics, 35 Kelinger Rd, Churchville, PA 18966-1033 (oren.livne@gmail.com, pinebirchanalytics.com; ORCID: 0000-0001-6700-483X).

and underlies graph partitioning [6]. Solving $L\phi = b$ is also a stepping stone to the Laplacian eigenproblem.

1.2. Related work. *Direct methods.* A Cholesky factorization $P^\top LP = \hat{L}\hat{L}^\top$ under a fill-reducing ordering (minimum-degree, nested dissection) is exact and, with a good ordering, very fast [7, 8]. But fill is governed by separator quality: $\mathcal{O}(n^{1.5})$ work on planar graphs [9], $\mathcal{O}(n^3)$ in general, which —on low-diameter social and citation graphs with no small separators—is memory-prohibitive.

Graph-theoretic iterative methods. Spielman and Teng [10] and successors [11] build spectral sparsifiers as preconditioners, achieving for any SDD system a provably near-linear $\mathcal{O}(m \text{ polylog } n \cdot \log(1/\varepsilon))$ bound. The hidden constants are large; the fast practical realization is the approximate-Cholesky (approxChol) solver [11, 12], recently made substantially more robust by Gao, Kyng, and Spielman (GKS) [13], and paralleled by randomized-Cholesky direct solvers such as RCHOL [14]. We benchmark against both approxChol variants as our strongest near-linear baselines (§4.3).

Algebraic multigrid. AMG [15, 16] builds a hierarchy of coarser operators from matrix entries alone. Classical AMG selects a coarse node subset; aggregation AMG [17, 18] groups fine nodes into aggregates. AMG excels on discretized PDEs but its complexity is hard to control on general graphs: coarse operators densify on scale-free hubs, and higher-order interpolation makes this worse. A line of work targets graph Laplacians specifically: Napov and Notay’s degree-aware rooted aggregation with aggregate-quality control [19] (the closest competitor, benchmarked in §4.4), matching-based multilevel preconditioners [20], and effective-resistance/diffusion-distance coarsening with least-squares interpolation from relaxed vectors [21]—the last closely related to LAMG’s affinity. Bootstrap AMG [22] and algebraic distance [23] likewise learn the coarsening from relaxed test vectors, at higher setup cost; a recent direction learns the prolongation itself with graph neural networks [24]. LAMG [1], by contrast, is parameter-free and training-free. Combinatorial Multigrid (CMG) [25] also targets general topologies; CMG and LAMG had comparable average speed but LAMG was markedly more robust. General-purpose aggregation packages (PyAMG [26], hypr/BoomerAMG [27]) are tuned for PDE discretizations rather than the irregular degree distributions of graph Laplacians. In the broadest recent comparison [13], CMG, BoomerAMG, and PETSc all failed to converge across the chimera and SDDM families, leaving the approximate-elimination solver AC (the `approxchol_lap2` variant) the only blackbox method robust across all of them; the original MATLAB LAMG was reported non-convergent there, whereas our independent re-derivation converges on those families (§4.6)—a failure specific to that evaluation, not the algorithm. We thus consider approxChol as our sole robust near-linear baseline, with CHOLMOD for the direct-method crossover. We also compare with the other AMG solvers for the specific graph classes they converge on.

1.3. Contribution. LAMG [1] is an accelerated caliber-1 aggregation AMG that attains $\mathcal{O}(m)$ on general graphs—more precisely $\mathcal{O}(m \log(1/\varepsilon))$ for an ε -accurate solution—by regulating the coarse-operator *energy* rather than aggregate size. We present **LAMG+**: an open re-derivation of the algorithm* (validated line-by-line against the original MATLAB reference) that we use to establish three things.

- *A fair, same-language adjudication against the modern near-linear champion* (§4.3). LAMG (2012) predates and was never compared with GKS’s approximate-Cholesky line [12, 13], now the dominant practical near-linear solver. Run in the same Julia process, the two are *complementary peers*: approxChol is faster on social/citation graphs, LAMG+ on

*Open-source Julia implementation, including all scripts to reproduce the experiments in this paper: <https://github.com/orenlivne/lamgplus>.

FE/structural matrices (the fastest robust solver there, $2.2\times$ over the robust AC variant on large graphs). The split is algorithmic (coarsening vs. fill), not a compiler artifact.

- *Linear scaling at scale* (§4.2). LAMG+ is empirically $\mathcal{O}(m)$ over 1,711 SuiteSparse graphs—100% converged, median 4 cycles, log–log slope 1.01 up to 2.4×10^8 nonzeros (twice larger than LAMG’s original demonstration) with no tuning.
- *Robustness, closing LAMG’s outlier tail* (§§3, 4.6). Where the original MATLAB LAMG was reported by [13] as non-convergent across the benchmark’s test families, we run the *unmodified* LAMG 2.2.1 on those families (regenerated from GKS’s released code) with the benchmark’s exact solver call and find it converges on every one, as does our re-derivation to the 10^{-8} tolerance (Table 4.6)—placing the reported failure in the evaluation, not the algorithm. Additionally, two lean, *local* refinements—a strength-of-connection veto and a selective caliber-2 interpolation—turn LAMG’s stalled grid-aligned-anisotropy rate (Asymptotic Convergence Factor (ACF) ≈ 0.99) into ≈ 0.11 at 1% median overhead, each acting only on the locally anisotropic nodes that need it.
- *An ablation study of LAMG’s load-bearing design choices* (§§2, 4.9). Several of LAMG’s setup choices—no aggregate-size cap, orphan nodes left as singleton seeds, a directional energy guard—are individually *load-bearing* yet undocumented as essential; we isolate them and quantify, by single-choice ablation, that replacing any one by a plausible correlate silently destroys convergence on a whole graph class (web ACF $0.002 \rightarrow 0.97$, for one). Here reproduction *is* discovery: knowing which heuristics carry the algorithm, and by how much, is what makes it transferable beyond the original implementation.

The open implementation makes all of the above reproducible; LAMG+ requires no parameter tuning.

2. The LAMG algorithm. We summarize LAMG [1] in enough detail to make this paper self-contained; the refinements of §3 presuppose it.

Multigrid in one paragraph. A relaxation such as Gauss–Seidel (GS) applied to $L\phi = b$ removes the high-residual (oscillatory) error in a few sweeps but stalls on the *algebraically smooth* error—the slow, low-energy modes. Multigrid corrects those on a coarser graph: it represents the fine smooth error $e \approx Pe^c$ by interpolation $P \in \mathbb{R}^{n \times n_c}$ from a coarse vector e^c , solves the Galerkin coarse system $L^c e^c = P^\top(b - L\tilde{\phi})$ with $L^c = P^\top LP$ and $\tilde{\phi}$ the approximation to ϕ after relaxation, and corrects $\tilde{\phi} \leftarrow \tilde{\phi} + Pe^c$. Applied recursively over $\ell = 1, \dots, L$ levels this is a multigrid cycle; the work is geometric in the coarsening ratio and therefore typically linear in m . LAMG uses two coarsening operators—elimination and aggregation—alternated until the coarsest level is small.

Relaxation. LAMG employs Gauss–Seidel, $\phi_u \leftarrow (b_u - \sum_{v \neq u} L_{uv}\phi_v)/L_{uu}$, an effective smoother for SPD systems that needs no damping parameter.

Low-degree elimination. Let $F \subseteq V$ be a maximal independent set of nodes of degree $\leq d_{\max} = 4$. Because F is independent, L_{FF} is diagonal and the Schur complement onto $C = V \setminus F$ is *exact*:

$$(2.1) \quad L^c = L_{CC} - L_{CF}L_{FF}^{-1}L_{FC}, \quad \phi_F = P_F\phi_C + L_{FF}^{-1}b_F, \quad P_F = -L_{FF}^{-1}L_{FC}.$$

Eliminating an F -node of degree ≤ 3 does not increase m ; the rare degree-4 node may add up to two edges. Elimination removes the effectively-one-dimensional part of the graph at zero approximation cost (up to fill), shrinking n before each aggregation step.

Caliber-1 aggregation. Each remaining node is assigned to exactly one aggregate, giving a piecewise-constant (“caliber-1”) interpolation with $P_{u,A(u)} = 1$. This keeps $L^c = P^\top LP$ as sparse as the graph—it cannot manufacture the dense coarse rows that sink classical AMG on hubs (hence the “lean” in LAMG). Aggregates are formed from K test vectors

$X = (x^{(1)}, \dots, x^{(K)})$, each being the result of ν GS sweeps on $Lx = 0$ from a random start: these are samples of the algebraically smooth subspace the coarse grid must capture, with no geometry assumed. (LAMG uses $K = 4$ at the finest level, raises K by one per coarser level up to 10, and smooths each by $\nu = 3$ sweeps [1, §4.1].) The proximity of an edge (u, v) is the *affinity*, the squared cosine of the two test-vector profiles [1, Eq. 3.3],

$$(2.2) \quad c_{uv} = \frac{(\sum_k x_u^{(k)} x_v^{(k)})^2}{(\sum_k (x_u^{(k)})^2)(\sum_k (x_v^{(k)})^2)} \in [0, 1].$$

A node u joins the seed s of maximal affinity that is *admissible* under the *energy-ratio guard*. With the nodal energy $E_u(x; y) = \frac{1}{2} L_{uu} y^2 - B_u(x) y + C_u(x)$, $B_u = \sum_v w_{uv} x_v$, $C_u = \frac{1}{2} \sum_v w_{uv} x_v^2$,

$$(2.3) \quad q_{u \leftarrow s} = \max_{1 \leq k \leq K} \frac{E_u(x^{(k)}; x_s^{(k)})}{\min_y E_u(x^{(k)}; y)} \leq Q, \quad Q = 2.5.$$

The numerator is u 's energy when caliber-1 *forces* $u = x_s$; the denominator is its relaxed (GS-optimal) energy $\min_y E_u = E_u(x; B_u/L_{uu})$. Their ratio is the factor by which lumping u into s 's aggregate inflates the local energy; capping it bounds the coarse-operator energy aggregate by aggregate, hence the two-level factor $\rho \lesssim 1 - 1/Q$ [1, §3.4]. Crucially the guard regulates *energy*, *not size*—there is *no* hard aggregate-size cap: the same Q admits a 100-node hub aggregate (neighbors agree on the slow modes, $q \approx 1$) yet forbids a 3-node mesh aggregate (neighbors disagree, q shoots up). Two companion conventions complete the rule and are equally load-bearing: a node that finds no admissible seed stays a singleton seed (an *orphan*), never force-absorbed; and the guard is *directional*—(2.3) scores u forced onto seed s , $q_{u \leftarrow s}$, and is not symmetrized in (u, s) . These are LAMG's choices [1, §3.4]; substituting a plausible correlate for any of them (a size cap, forced absorption, a symmetric guard) silently over- or under-coarsens some graph class, so §3 discusses where LAMG+ deviates from LAMG. The guard dominates setup cost and is evaluated in $\mathcal{O}(m)$ via three bulk-precomputed moments per node (App. B).

Cycle and acceleration. The solve runs a γ -cycle whose per-level index LAMG fixes by a work rule [1, Eq. 4.1]: elimination levels use $\gamma_\ell = 1$, and an aggregation level uses

$$(2.4) \quad \gamma_\ell = \begin{cases} 1.5, & |E_\ell| > 0.1 |E| \text{ (finer levels),} \\ \min\{2, 0.7/\tau_\ell\}, & \text{otherwise (coarser levels),} \end{cases}$$

where $\tau_\ell = |E_{\ell+1}|/|E_\ell|$ is the level- ℓ coarsening ratio. The cap bounds per-cycle work geometrically, guaranteeing $\mathcal{O}(m)$. The coarse-level branch raises γ_ℓ toward 2 where coarsening is aggressive but lowers it toward 1 where coarsening is weak ($\tau_\ell \rightarrow 1$); LAMG+ keeps the same rule with a slightly looser work cap. After each cycle the finest iterate is improved by a min-residual (Krylov) recombination over a short history of κ saved iterates [28, §7.8.2]: $\alpha = \arg \min_\alpha \|r - LE\alpha\|$ with $r = b - L\tilde{\phi}$, $E = [\phi_1 - \tilde{\phi}, \dots, \phi_\kappa - \tilde{\phi}]$, then $\tilde{\phi} \leftarrow \tilde{\phi} + E\alpha$. A level's history accumulates across the cycle's γ revisits; the finest persists across cycles.

3. Two refinements beyond LAMG. Reproducing LAMG (§2)—no size cap, orphan singletons, a directional guard, the cycle-index calibration—is a prerequisite, not a contribution. LAMG+ departs from LAMG (Table 3.1) by (1) avoiding aggregation across a weak edge, via a strength-of-connection veto (§3.1), and (2) judiciously raising the interpolation caliber from 1 to 2 (§3.2); both target (local) grid-aligned anisotropy—the bounded-but-slow outlier tail LAMG reported [1, §§5.2, 6.1]. Together they turn a stalled ACF ≈ 0.99 into ≈ 0.11 on the

worst anisotropic grids while leaving every other class untouched. Our contribution is their lean, parameter-free realization and at-scale evidence (§4). Both refinements are *local*—the veto a per-edge test, caliber-2 a per-node test, each firing only where the matrix is locally anisotropic and inert ($\approx 0\%$ of nodes) elsewhere—so refinement follows where the solver is locally slow, never a global average: a graph that is mostly isotropic with a small anisotropic region defeats any global trigger yet is handled correctly, node by node (§3.3). Because the two act at different stages (which nodes aggregate, how they interpolate), they are independent and individually ablatable (§4.9).

TABLE 3.1

The two places LAMG+ departs from LAMG, both targeting grid-aligned anisotropy. Reproduced setup choices (no size cap, orphan singletons, directional guard, non-decaying γ , fixed $K=4$) are discussed in §§2, 3.3.

Refinement	LAMG	LAMG+	Effect
Aggregation	affinity + energy guard only	+ strength-of-connection veto: no merge across a weak edge (eq. (3.1))	removes the rare affinity-inversion stall, ACF $0.99 \rightarrow 0.46$
Interpolation caliber	strictly caliber-1	+ selective caliber-2 on locally-1-D nodes	lowers the $q \approx 2$ ceiling, ACF $0.46 \rightarrow 0.11$

3.1. A strength-of-connection veto. Affinity (2.2) ranks a node’s neighbors by how well their relaxed test-vector profiles align, and aggregation merges a node with its highest-affinity admissible neighbor. On a strongly grid-aligned anisotropic grid the test vectors are nearly flat along the weak direction, so the K -sample affinity of a weak (weight- ϵ) edge is computed from almost-identical profiles and can, on an unlucky sample, *exceed* that of the genuinely strong (weight-1) edge—we measure 0.97 versus 0.89 at one interior node of a 256^2 , $\epsilon = 10^{-4}$ grid. Neither the affinity nor the energy guard Q catches this: both are computed from the *same* test vectors and are fooled identically. The *matrix* is not—the conductance ratio there is 10^4 . A single such inversion creates one cross-weak-direction merge, a localised residual hot-spot (peak-to-mean $> 10^3$), and a stalled rate (ACF ≈ 0.99) that no amount of γ or recombination repairs, and that caliber-2 alone cannot reach past either.

The veto is the classical Ruge–Stüben strength-of-connection criterion [16] applied to aggregation: never aggregate across an edge whose weight is below a fraction τ_{soc} of the node’s strongest incident weight,

$$(3.1) \quad \text{skip edge } (u, v) \text{ in aggregation if } |w_{uv}| < \tau_{\text{soc}} \max_{t \sim u} |w_{ut}|, \quad \tau_{\text{soc}} = 0.05.$$

It costs $\mathcal{O}(1)$ per edge, changes no asymptotics, and is inert wherever a node’s incident weights are comparable. With it, the 256^2 and 384^2 grids drop from ACF ≈ 0.99 to 0.11 across all sizes, seeds, and $\epsilon \leq 10^{-1}$ (worst 0.117 over 24 cases); on a 42-graph random SuiteSparse sample and the social/FE giants to 1.5×10^7 edges it is neutral (worst ACF ratio 1.22, no regression). The defect is rare—it needs a near-degenerate weak direction *and* an unlucky sample—which is why stock LAMG mostly avoids it and reports anisotropy as bounded-but-slow rather than stalled; the veto removes the residual tail and, crucially, is what lets the caliber-2 fit of §3.2 act on a clean aggregation.

3.2. Selective caliber-2 on one-dimensional chains. The veto fixes *which* nodes aggregate; caliber-2 changes *how* they interpolate—the place LAMG’s strict caliber-1 design

TABLE 3.2

Caliber-1 interpolation vs. caliber-2 (ACF, with cycles to $\varepsilon = 10^{-8}$ in parentheses). “aniso. nodes” is the fraction of locally one-dimensional nodes. Non-anisotropic graphs are unaffected (§4.9).

graph	aniso. nodes	caliber-1	caliber-2 (default)
aniso 256^2 , $\epsilon=10^{-4}$	99%	0.56 (13)	0.16 (7)
aniso 384^2 , $\epsilon=10^{-4}$	99%	0.83 (33)	0.11 (6)
isotropic $256^2 + 64^2$ aniso. patch	6%	0.31 (8)	0.02 (5)
isotropic $512^2 + 64^2$ aniso. patch	1.5%	0.24 (7)	0.02 (5)
isotropic grid 256^2	0%	0.016 (5)	0.015 (5)
bodyy5 (FE stiffness)	0%	0.020 (5)	0.019 (5)

leaves a bounded but non-textbook rate. On grid-aligned anisotropy the affinity correctly semicoarsens along the strong direction (the aggregates are well placed), but the caliber-1 piecewise-constant interpolation across each 2:1 aggregate recovers only half of the smooth-in-strong mode’s amplitude, inflating the energy ratio to $q \approx 2$ and pinning the two-level factor at $\rho \approx 0.5$ (App. A). Where a fine node’s *strong* edges (i.e., those not vetoed by (3.1)) reach exactly two coarse aggregates—a locally one-dimensional neighborhood—LAMG+ adds a second, caliber-2 parent with a single fitted weight w (App. A derives it). Measured per level with an exact coarse solve (so the γ schedule cannot mask it), the two-level factor on the 256^2 , $\epsilon = 10^{-4}$ grid falls from 0.46 (caliber-1, veto on) to 0.28 at every aggregation level; with γ and recombination the full-cycle rate reaches ≈ 0.11 . The gate is *self-targeting*: $\approx 0\%$ of nodes upgrade on isotropic graphs, so over a 1,267-graph SuiteSparse sample it rescues the anisotropic minority (NACA0015, epb3, darcy from 100 cycles to 5–8; the thermomech/thermal stiffness matrices from non-convergence to ≈ 13) while leaving the easy majority’s solve essentially unchanged, at $\approx 1\%$ median wall-clock. The two refinements are complementary: caliber-2 without the veto still stalls at 0.99 on the worst grids, the veto without caliber-2 reaches only the caliber-1 ceiling ≈ 0.46 ; together they give 0.11 (§4.9).

3.3. A single, local configuration. LAMG+ runs one configuration on every graph: a cheap $K=4$ affinity with both refinements always on. The interpolation caliber is raised from 1 to 2 wherever a node is locally one-dimensional (§3.2) and left at caliber-1 everywhere else, so the extra cost falls only on the anisotropic nodes. Four test vectors are an ample sample for the single least-squares weight each upgraded node fits, so we simply fix $K=4$ rather than accumulate test vectors with depth as the original LAMG does (raising K helps reduce the probability of a spurious affinity, but would significantly pay off only together with more relaxation sweeps ν to increase the interpolation caliber beyond 2, which we avoid in this lean AMG framework).

Because the caliber decision is taken per node, it needs no global control, making it robust. A graph that is mostly isotropic with a small anisotropic region is handled correctly: the upgrade fires exactly on that region. In a 64^2 anisotropic patch embedded in a 256^2 or 512^2 isotropic grid, only 6% and 1.5% of the nodes trigger, respectively (Table 3.2). The caliber-1 ACF ≈ 0.3 improves to a textbook multigrid efficiency ACF ≈ 0.02 for caliber-2.

4. Numerical results. LAMG+ is implemented in Julia, and all reported times are wall-clock, measured on an Apple M5 Pro (18-core, 48 GB RAM); the competitor solvers of §4.3 run in the same Julia process under identical compilation, so the wall-clock comparisons are language-fair. Throughout this section, every solve uses the right-hand side $b = L x_{\text{true}}$ with x_{true} a standard-normal random vector projected to zero mean, and runs to $\varepsilon = \|b - Lx\|/\|b\| \leq 10^{-8}$ unless a different ε is stated.

TABLE 4.1

ACF on identical (L, b) ; “cyc” is cycles to $\varepsilon = 10^{-8}$. “LAMG+” and “ref.” are the LAMG+ and MATLAB-reference asymptotic convergence factors. Entries $< 10^{-4}$ solve in 1–2 cycles—below the resolution at which an asymptotic factor is meaningful. Class abbreviations: AS = autonomous-systems, P2P = peer-to-peer.

Graph	Class	n	m	LAMG+	cyc	ref.
grid128	structured	16 384	32 512	0.022	5	0.21
bodyy5	FE stiff.	18 589	55 132	0.019	5	0.27
wb-cs-stanford	web	7 008	17 054	0.0009	3	0.06
ca-GrQc	collab.	4 158	13 422	0.0008	3	0.07
ca-HepTh	collab.	8 638	24 806	0.0040	4	0.08
Oregon-1	AS	11 174	23 409	$< 10^{-4}$	2	0.005
as-caida	AS	26 475	53 381	$< 10^{-4}$	2	0.008
p2p-Gnutella08	P2P	6 262	9 719	0.0010	3	0.03
email-Enron	social	33 696	180 811	0.0011	3	0.10
ak2010	road	42 381	102 091	0.0066	4	0.12

Test corpus. All experiments draw from one fixed corpus, the union of two sources: (i) the complete 1,711-graph SuiteSparse collection used as the LAMG test set ($100 \leq n \leq 1.2 \times 10^7$), which drives the scaling study (§4.2) and, restricted to its $m > 10^6$ members, the head-to-head competition (§4.3); (ii) instances from the AC robustness benchmark [13], generated by its own generators—all 30 `chimera` and `weighted-chimera` graphs and all 9 uniform/anisotropic/high-contrast Poisson grids—plus our synthesized Kyng–Sachdeva classes (Erdős–Rényi, Barabási–Albert, stochastic-block, k -regular, power-law-cluster), used for robustness (§4.6). Every solver runs on the same exported (L, b) ; all per-graph setup and solve times are recorded and released with the code [29] so the tables can be regenerated. The largest graph solved end-to-end is `com-Orkut` (1.2×10^8 edges, 2.4×10^8 nonzeros, 2 cycles).

4.1. Convergence against the reference. Table 4.1 reports the asymptotic convergence factor (ACF, the geometric mean of the last five per-cycle residual reductions) on identical exported systems (L, b) . LAMG+ matches or beats the MATLAB reference on every class.

4.2. Linear scaling. Figure 4.1 plots LAMG+ per-edge setup and solve time versus m for all 1,711 SuiteSparse graphs ($100 \leq n \leq 1.2 \times 10^7$). Every one converges to $\varepsilon = 10^{-8}$, in a median of 4 cycles. The log–log slopes of per-edge time are $+0.005$ (setup) and -0.002 (solve): per-edge cost is independent of m , supporting the empirical $\mathcal{O}(m)$ statement (equivalently, total setup+solve time scales as $m^{1.01}$). The medians are $t_{\text{setup}}/m \approx 0.24 \mu\text{s/edge}$ (IQR $[0.18, 0.34]$) and $t_{\text{solve}}/m \approx 0.19 \mu\text{s/edge}$ (IQR $[0.10, 0.43]$). These low, flat constants reflect an implementation tuned to the algorithm: a self-gated reverse-Cuthill–McKee reordering [35] for cache locality, single-pass Gustavson-fused Schur extraction (each coarse operator built in one column scan with no matrix-multiply temporary), allocation-free elimination into preallocated buffers, and a monomorphic, branch-free, residual-maintaining in-place Gauss–Seidel kernel. Each of these features does not change the result but lowers the hidden constant. The largest per-edge costs are the strongly anisotropic grids; cf. App. A.

4.3. Comparison: complementary strengths by class. We benchmark setup+solve to $\varepsilon = 10^{-8}$ against both approximate-Cholesky variants: the original, faster `approxchol_lap` (“approxChol”) [12] and the robust `approxchol_lap2` (“AC”) of [13], and note sparse Cholesky (CHOLMOD) [7] for context.

All run in the same Julia process under identical compilation, so any difference is algorithmic (coarsening vs. fill), not a compiler artifact; we report time per nonzero ($\mu\text{s}/\text{nnz}$,

LAMG+ linear scaling — 1,711 real-world graph Laplacians ($m = 52..982,542$)




FIG. 4.1. LAMG+ per-edge time vs. number of edges m (log-log) over 1,711 SuiteSparse graph Laplacians: (a) setup, (b) solve, (c) total. Grey points are individual graphs; the red dashed line is the median, the red solid line the log-log regression fit (slope shown per panel). A near-zero slope means per-edge cost is independent of m — the empirical $\mathcal{O}(m)$ statement; the regression of total setup+solve time gives the exponent $m^{1.01}$.

TABLE 4.2

Setup+solve time per nonzero ($\mu\text{s}/\text{nnz}$) to $\varepsilon = 10^{-8}$; bold = fastest. “approxChol” is the fast `approxchol_lap`; “AC” the robust `approxchol_lap2`. Sparse Cholesky is omitted: it is fastest on the web/road graphs but cannot factor the FE/structural and dense-citation matrices (> 21 GB fill on `flickr`). The bottom block summarizes LAMG+ against each competitor over all 201 graphs with $m > 10^6$ ($\text{nnz}(L) \leq 6 \times 10^7$): the geometric-mean speedup is the geometric mean of (competitor time)/(LAMG+ time), so > 1 means LAMG+ is faster; the win rate is the fraction of the 201 graphs on which LAMG+ is faster; each reported at $\varepsilon = 10^{-8}/10^{-4}$. All times are serial (contention-free).

Graph	family	edges	LAMG+	approxChol	AC
web-Stanford	web	1.0 M	0.26	0.19	0.30
web-Google	web	2.6 M	0.48	0.25	0.43
roadNet-PA	road	1.5 M	0.72	0.26	0.34
flickr	social	4.9 M	0.55	0.26	0.62
coPapersDBLP	citation	15.2 M	0.31	0.23	0.46
bmwcra_1	FE	5.2 M	0.08	0.28	0.57
crankseg_1	structural	5.3 M	0.08	0.19	0.65
troll	structural	5.9 M	0.10	0.24	0.44
pwtck	structural	5.7 M	0.11	0.21	0.40
all $m > 10^6$ graphs (201), LAMG+ vs. each ($10^{-8}/10^{-4}$):					
geometric-mean speedup				0.99/1.20×	2.20/2.91×
win rate				56/64%	76/84%

the worst-case metric of [13]), each a warm timed run (a warm-up call first excludes Julia’s one-time JIT compilation), measured serially so the per-graph rows and the aggregate below are contention-free and mutually consistent.

The finding is a clean *class split*, not a single winner (Table 4.2), but at root it is one mechanism—a *fill contest*. These graphs all converge in only 3–4 cycles, so neither solver is paced by iteration count; the winner is whichever builds the *leaner hierarchy*, because the fill a solver creates sets both its setup cost and the cost of every cycle. Aggregation and sampled elimination have *opposite* fill profiles. On low-diameter social, web, and citation graphs and on the planar road network, approxChol wins: their low mean degree keeps its sampled elimination cheap, while the absence of clean separators inflates LAMG+’s coarse operators—and with them the cost of every cycle. On finite-element and structural matrices the profiles reverse—the high per-node coupling makes elimination fill in, while LAMG+’s aggre-

TABLE 4.3

Multi-solver comparison by graph class: mean/worst-case setup+solve time per nonzero ($\mu\text{s}/\text{nnz}$) to $\varepsilon = 10^{-8}$, over a class-stratified corpus (real SuiteSparse/SNAP classes and all of the approximate-Cholesky benchmark’s synthetic families, including the SPE10 reservoir). Bold = fastest mean; “DNF” = all instances failed to converge. A per-graph budget of $50\times$ the fastest solver (capped at 90 s) bounds the slow cases. $\text{apxChol} = \text{approxchol_lap}$; $\text{PETSc} = \text{PETSc GAMG}$.

class	LAMG+	apxChol	AC	Boomer	pyAMG	CMG	PETSc
FE/structural	0.12/0.15	0.28/0.54	0.53/0.73	0.18/0.35	0.10 /0.13	0.18/0.21 ²	0.12/0.20
mesh/grid	0.84/1.24	0.24/0.32	0.55/0.74	0.18 /0.27	0.38/0.54	1.46/1.46 ²	0.23/0.33
social/citation	0.47/0.69	0.25 /0.37	0.71/0.94	1.53/2.50	4.60/7.81 ³	0.50/0.73	3.10/3.14 ⁴
web	0.37/0.41	0.29 /0.38	0.39/0.51	1.43/2.59	5.45/5.45 ³	1.27/2.05 ¹	2.76/3.59 ²
road	0.68/0.85	0.39/0.50	0.53/0.65	0.35 /0.39	1.38/1.76	1.24/1.24 ³	0.57/0.75
optimization	0.15/0.18	0.26/0.36	0.88/1.03	4.07/7.02	0.13 /0.13	0.23/0.33	2.04/3.47
chimera	0.40/0.91	0.14 /0.19	0.56/1.20	0.48/0.67	1.13/1.74 ¹	0.43/0.78	0.91/1.72
wtd-chimera	0.45/1.34	0.23 /0.70	0.39/0.56	0.46/0.63	2.04/3.25 ¹	0.33/0.76	0.96/1.86
sddm-chimera	0.30/0.51	0.10 /0.13	0.21/0.32	0.44/0.57	2.21/2.80 ¹	0.34/0.47	1.74/3.26
grid-3D	1.30/1.35	0.32/0.37	0.70/0.72	0.19 /0.19	0.37/0.41	DNF	0.25/0.25
aniso-grid	0.61/0.68	0.15/0.16	0.38/0.41	0.15 /0.15	2.38/2.38 ²	1.22/1.48 ¹	1.01/1.49
star	0.05/0.05	0.03/0.03	0.24/0.40	0.02 /0.02	0.03/0.03	0.04/0.06	0.04/0.04
SPE/reservoir	1.09/1.22	0.29 /0.30	0.56/0.59	0.29/0.31	2.72/2.95	DNF	0.90/1.29

Superscript k : k instances in the class exceeded the 300-cycle limit; shown values are medians over converged instances only.

gation collapses those dense neighborhoods at $\mathcal{O}(1)$ cost on a cleanly separating mesh—and so does the order: **LAMG+ is the fastest of the three** (geometric-mean $1.76\times$ over approxChol across the 34 FE/structural graphs, winning 91% of them; `bmwcra_1`, `crankseg_1`, `troll`, `pwtk` run $1.9\text{--}3.5\times$ over it), the leanest build of the three, and the fixed cheap $K=4$ calibration (§3.3) pays off. Aggregated over all 201 large graphs ($m > 10^6$), the two are *near-peers*: LAMG+ wins 56% of head-to-heads against approxChol (geometric-mean $0.99\times$ —approxChol’s social/web wins are by similar factors to LAMG+’s FE/structural wins), and 76% against the robust AC, where it is $2.2\times$ faster in geometric mean. (At the lower 10^{-4} accuracy many applications need, the balance shifts toward LAMG+; cf. §4.7.) Sparse Cholesky is fastest where good separators exist (web crawls, road) but cannot factor the largest poorly-separable matrices at all.

Table 4.3 widens the comparison to all seven solvers across 13 graph classes, including all of the AC benchmark’s synthetic families. The result is a clean separation by design intent. On the structured classes they were built for—FE/structural and grids—the four algebraic-multigrid solvers (LAMG+, pyAMG, BoomerAMG, PETSc GAMG) are all fast and beat sparse Cholesky. On the low-diameter classes (social, citation, web) BoomerAMG [27] and PETSc GAMG slow by up to an order of magnitude, pyAMG [26] stops converging, and approxChol is fastest. CMG [25] does not converge on the grids, on SPE, or on many roads. BoomerAMG converges everywhere but is catastrophic ($7\ \mu\text{s}/\text{nnz}$) on irregular optimization graphs; pyAMG, CMG, and PETSc GAMG additionally fail to converge on 11, 15, and 6 of the instances, respectively. Across all classes only three solvers are both convergent everywhere and never catastrophic (worst case under $1.4\ \mu\text{s}/\text{nnz}$): **LAMG+** (≤ 1.35), approxChol (≤ 0.70), and AC (≤ 1.20)—precisely the robust near-linear solvers this paper compares.

Memory. The same split holds in space (Table 4.4). We compare memory only against the robust solvers (approxChol, AC): the non-robust competitors store their hierarchies in C and are not directly measurable in bytes/nnz. On large FE/structural matrices LAMG+’s no-fill multilevel hierarchy is the most frugal of the three (21–30 bytes/nnz vs. 33–55 for the factorizations); on web/social graphs it is the most expensive. Across every graph the robust AC costs more memory than approxChol, so AC’s guarantee is paid in both time and space.

TABLE 4.4

Memory usage (bytes per nonzero of L) of the three robust solvers. LAMG+ is the most frugal on FE/structural matrices, heavier on web/social; AC always exceeds `approxchol_lap`. Non-robust solvers (BoomerAMG, CMG, pyAMG, PETSc GAMG) store hierarchies in C and are not included.

Graph	LAMG+	approxChol	AC
troll, pwtk (structural)	23	34	44
bmwcra_1 (FE)	22	35	47
bone010 (FE, 2.3×10^7)	30	43	55
flickr (social)	80	40	53
coPapersDBLP (citation)	49	32	32
web-Stanford (web)	62	32	32

The table reports the *persistent* hierarchy; the transient setup peak adds only a single level’s K test vectors ($K \cdot n$ floats), discarded once that level is coarsened rather than accumulated across levels, and the fixed $K=4$ default halves this peak versus LAMG’s $K=8$ accumulation, which governs whether the largest graphs build within a fixed memory budget.

4.4. The closest competitor: degree-aware aggregation. The graph-Laplacian solver most closely related to LAMG is the degree-aware rooted aggregation with quality control (DRA-QC) of Napov and Notay [19]: like LAMG it eliminates low-degree vertices and coarsens by aggregation, but it controls aggregate quality through a provable two-grid condition-number bound rather than an energy guard, and is reported $2\text{--}5\times$ faster than the original LAMG on a large real-world test set. No implementation of the DRA-QC variant is publicly distributed and we were unable to obtain the authors’ code; we therefore reimplemented it faithfully—validated to reproduce its published operator complexities and $15\text{--}30$ -iteration range—and emphasize that *all DRA-QC figures below are from this reimplementation*, not the original. We ran all three robust near-linear solvers under identical build-once conditions over the full SuiteSparse corpus of §4.2.

All three are robust—each converges on $> 99.7\%$ of the corpus (Table 4.5)—but LAMG+ is the only one that converges on *every* graph, and with the tightest tail (at most 27 cycles, versus 77 for approxChol and 277 for DRA-QC); the stragglers are anisotropy/Darcy/thermal matrices for DRA-QC and a single chemical-process matrix for approxChol. Measured in work per edge (the identical operation count for LAMG+ and our DRA-QC reimplementations), the two are *complementary*: DRA-QC’s leaner cycle does *less* total work than LAMG+ on the typical graph (median ratio 0.68), but over $3\times$ on the 11% of graphs with grid-aligned anisotropy or high contrast (up to $55\times$), where aggregation-only coarsening lacks LAMG+’s caliber-2 fix, while LAMG+’s worst-case work stays bounded. DRA-QC is thus a strong competitor with a complementary cost profile—robust on real-world graphs with a proven bound—but without LAMG+’s uniform convergence and bounded tail.

4.5. An a-priori selection rule. We construct a simple predictor of a LAMG+ win from a single statistic computable before either solver runs: the mean degree $\bar{d} = 2m/n$. Treating “LAMG+ wins iff $\bar{d} \geq 30$ ” as a binary classifier over the 201 large graphs, its accuracy peaks at 91% near $\bar{d} \approx 30$ (Figure 4.2), and the point-biserial correlation of $\log \bar{d}$ with a LAMG+ win is $r = +0.68$. The mechanism is two cost axes keyed to properties that anti-correlate across this corpus: approximate-Cholesky’s fill grows with elimination-clique size, hence with degree, while LAMG+’s operator complexity grows with poor separability (low-diameter, small-world structure)—and here the well-separated graphs are the high-degree ones, so $\bar{d}$ proxies both. High $\bar{d}$ favors LAMG+ and low $\bar{d}$ favors sampled elimination; the residual 9% are exactly the

TABLE 4.5

Robustness over the full SuiteSparse corpus (build the hierarchy once, solve to 10^{-8}): fraction of graphs converged and the per-graph iteration distribution. LAMG+ caliber-2 is the only solver convergent on every graph and has the tightest tail. “it” = cycles (LAMG+), PCG iterations (approxChol), or K-cycle iterations (DRA-QC, our faithful reimplementation, as no public code exists).

Solver	converged	median it	p99 it	max it	#fail
LAMG+ (caliber-2)	100.0%	4	9	27	0
approxChol [13]	99.9%	15	37	77	1
DRA-QC [19]	99.8%	13	89	277	4




FIG. 4.2. Degree-threshold predictor of a LAMG+ win: accuracy of the rule “LAMG+ wins iff $\bar{d} > t$ ” over the 201 large ($m > 10^6$) graphs as a function of the threshold t , at both solve accuracies. At $\varepsilon = 10^{-8}$ (red, solid) accuracy peaks at 91.0% near $\bar{d} \approx 30$; at the looser $\varepsilon = 10^{-4}$ (blue, dashed) LAMG+’s higher win rate (64% vs. 56%; §4.7) shifts the optimal threshold left, peaking at 90.5% near $\bar{d} \approx 13$. Both peaks are marked. The near-balanced $\varepsilon = 10^{-8}$ classes (112 LAMG+ vs. 89 approxChol wins) give a 56% majority-class baseline, and the point-biserial correlation of $\log \bar{d}$ with a LAMG+ win is $r = +0.68$ ($+0.72$ at $\varepsilon = 10^{-4}$). The mean degree thus predicts the faster solver a priori at either accuracy.

graphs where the two axes disagree—high-degree but small-world, or low-degree but cleanly separable.

Does a richer model help? We tested a learned selector to try to explain the residual 9%: a gradient-boosted decision tree (a histogram model of the LightGBM class [37]) trained on 13 further a-priori features of the degree distribution—median, maximum, 90th/99th percentiles, standard deviation, coefficient of variation, skewness, hub ratio $\max / \bar{d}$, and the leaf- and high-degree fractions, all computable before any solve. Under repeated stratified 5-fold cross-validation (20 repeats; a single 80/20 split is too small to be stable at $n=201$) it attains $89 \pm 4\%$ —no better than the one-line $\bar{d}$ threshold (91%); permutation importance concentrates almost entirely on $\bar{d}$, every other feature contributing ≤ 0.01 . The degree distribution thus carries no a-priori signal beyond its mean; the residual 9% is governed by separability, a global property no cheap local statistic exposes, so the principled one-statistic rule is also the near-optimal one. A black-box library can apply it directly—dispatch to LAMG+ when $\bar{d} \gtrsim 30$ and to approximate-Cholesky otherwise—to pick the faster solver per input at no measurement cost.

4.6. Robustness across the GKS test families. The crossover above is a speed result; robustness—converging at all—is a separate axis, and the one on which graph-Laplacian solvers most often differ. Gao, Kyng, and Spielman (GKS) [13] report AC as the only solver in their study (alongside CMG, hypr/BoomerAMG, PETSc, and ICC) to converge across

TABLE 4.6

The original MATLAB LAMG 2.2.1 vs. LAMG+ on the families of Gao-Kyng-Spielman (GKS) approximate-Cholesky benchmark [13], same graph and right-hand side. Both LAMG and LAMG+ converge for all instances. LAMG+ reaches a strict 10^{-8} residual in as few or fewer cycles. “cyc” = multigrid cycles; relres = $\|L\phi - b\|/\|b\|$ at termination. (LAMG 2.2.1’s `errorReductionTol` test targets the error norm, so its residual lands within a small factor of 10^{-8} , e.g., 3.2×10^{-8} for uniform-weight chimera.)

Family	n	m	LAMG 2.2.1		LAMG+	
			cyc	relres	cyc	relres
chimera (unweighted)	50,000	173,059	10	1.8×10^{-9}	5	8.5×10^{-9}
chimera (uniform wt)	50,000	99,552	9	3.2×10^{-8}	5	4.9×10^{-10}
chimera (weighted)	50,000	76,068	8	1.0×10^{-9}	4	1.6×10^{-9}
chimera (SDDM)	50,000	173,059	10	2.7×10^{-9}	5	6.6×10^{-9}
Sachdeva star	10,001	495,100	4	2.7×10^{-10}	3	4.6×10^{-10}
grid 2D (uniform)	65,536	130,560	11	3.8×10^{-9}	5	6.4×10^{-10}
grid 2D (anisotropic)	65,536	130,560	19	2.3×10^{-9}	6	7.8×10^{-9}
grid 2D (high-contrast)	65,536	130,560	10	6.0×10^{-9}	4	6.6×10^{-9}
grid 3D (uniform)	64,000	187,200	10	8.5×10^{-9}	9	6.1×10^{-9}
grid 3D (anisotropic)	64,000	187,200	16	8.1×10^{-9}	11	2.2×10^{-9}
grid 3D (high-contrast)	64,000	187,200	12	7.1×10^{-9}	12	3.7×10^{-9}
SPE10 (reservoir)	264,000	773,200	11	4.8×10^{-9}	8	6.9×10^{-9}

their full family set—random chimera graphs and randomly-weighted variants, SDDM “chimeras”, adversarial “Sachdeva stars”, an SPE benchmark, and uniform/anisotropic/high-contrast Poisson grids. That robustness is precisely why we take AC as the near-linear baseline worth a same-machine comparison.

LAMG+ meets the same bar. With the standard right-hand side ($b = L x_{\text{true}}$, zero-mean), it converges at both $\varepsilon = 10^{-8}$ and $\varepsilon = 10^{-4}$ on every one of the benchmark’s families generated by GKS’s own `Laplacians.jl`: chimeras and their weighted, uniform, and semi-weighted variants (96 graphs, $n = 10^4$ to 10^6), SDDM chimeras (12), 3-D Poisson grids with uniform, anisotropic, and high-contrast coefficients (12), and Sachdeva-style stars (3). It also converges on the benchmark’s one *external* family, the SPE fluid dataset: we reconstruct the SPE10 reservoir model [30] directly from its public permeability field, discretized by a two-point flux approximation into high-contrast (permeability ratio $\sim 10^7$), anisotropic graph Laplacians at three sizes ($n = 2.6 \times 10^5$ to 1.1×10^6), each solved in 5 cycles. Extending *past* the scale at which AC was benchmarked, it further solves 41 large SuiteSparse graphs up to `com-Orkut` (1.2×10^8 edges, 2.4×10^8 nonzeros). Together with the 1,711-graph scaling set (§4.2), every instance converges, in a median of 4 cycles. The original MATLAB LAMG was reported non-convergent on all matrices of this set [13]. This is not a property of the algorithm or of its original code: running the *unmodified* LAMG 2.2.1—under GNU Octave (compatibility shims only; the multigrid algorithm, setup, and MEX kernels unchanged), with [13]’s exact solver call (`setup('laplacian', ., errorReductionTol=10-8, default cycle cap)`)—on the same families converges on every one (Table 4.6), as does LAMG+, in as few or fewer cycles and to a strict 10^{-8} residual. The reported failure is therefore an artifact of the original evaluation, not the solver. (We could not reproduce the original evaluation’s failure mechanism.)

Our first-hand results corroborate the central finding of [13]: BoomerAMG, CMG, pyAMG, and PETSc GAMG each fail to converge or slow by an order of magnitude on chimera, SDDM-chimera, or adversarial-star instances (Table 4.3); the present comparison adds LAMG+ as the only other solver, alongside AC, that is robust across all twelve families. On these well-structured families `approxChol` is faster—bounded degrees and separators suit its sampled elimination—while LAMG+’s advantage is on large poorly-separable graphs.

TABLE 4.7

When each choice is necessary, and its effect when wrong (negligible elsewhere). “Stresses” is the graph class on which the choice is load-bearing. Top: LAMG choices LAMG+ reproduces (§§2, 3.3); bottom: the two anisotropy refinements.

Choice	Stresses	Effect if wrong
no size cap	scale-free / web / social	web ACF 0.002 $\rightarrow$ 0.97 (dense fill, stall)
orphan singletons	anisotropic FE	bodyy5 diverges ($\rightarrow$ 0.99)
directional guard	structured	over-aggregation; rarely fatal alone
strength-of-connection veto	grid-aligned anisotropy	rare affinity-inversion stall ($\rightarrow$ 0.99)
caliber-2 (default on)	grid-aligned anisotropy	ceiling 0.46 $\rightarrow$ 0.11; 100-cyc near-failures $\rightarrow$ 5–14

4.7. Lower accuracy shifts the balance toward LAMG+. Many applications need only a few digits. At $\varepsilon = 10^{-4}$ the LAMG+ solve collapses to a *single* cycle. This helps LAMG+ across the board (Table 4.2, bottom block) for a concrete reason: LAMG+ spends the larger share of its time in the *solve* (a median 68% of its total on the low-diameter class, versus 37% for approximate-Cholesky)—a few cycles over its higher-complexity hierarchy outweigh the lean PCG—so shedding it pulls every graph toward its setup-only ratio, where LAMG+ stands relatively better, even on the graphs it still loses. Against the approxChol the aggregate geometric mean rises from $0.99\times$ at $\varepsilon = 10^{-8}$ (parity) to $1.20\times$ at $\varepsilon = 10^{-4}$, and its win rate over the 201 large graphs rises from 56% to 64%. Against the robust AC the lead widens from $2.20\times$ to $2.91\times$ (76% $\rightarrow$ 84% of head-to-heads). The gain is broad but does not overturn the low-diameter classes—there it only narrows LAMG+’s deficit; the decisive flips are the moderate-degree graphs near the crossover. The degree-threshold predictor tracks this shift: its optimal cutoff reduces from $\bar{d} \approx 30$ at $\varepsilon = 10^{-8}$ to ≈ 13 at $\varepsilon = 10^{-4}$ (Figure 4.2).

4.8. Anisotropy: the veto and the caliber-2 extension. Grid-aligned anisotropy is baseline (caliber-1) LAMG+’s worst class. On a 64×64 grid with strong weight 1 and weak weight ϵ the full solver converges at a bounded but non-textbook ACF $\approx 0.27\text{--}0.39$ for $\epsilon \leq 10^{-1}$ (vs. 0.01 isotropic); the caliber-2 gate (§3.2) recovers the ideal rate $\text{ACF} = 0.007/0.05/0.10$ at $\epsilon = 10^{-1}/10^{-2}/10^{-4}$, at *lower* operator complexity (sharper interpolation buys fewer levels). On larger grids (256^2 , 384^2) a second, rarer failure appears: the affinity inversion of §3.1 produces a single bad merge that stalls even caliber-2 at $\text{ACF} \approx 0.99$. The strength-of-connection veto removes it, and the two refinements together hold $\text{ACF} \approx 0.11$ across all sizes, seeds, and $\epsilon \leq 10^{-1}$ (worst 0.117 over 24 cases). On the SuiteSparse stiffness matrices the pair converts 100-cycle near-failures into 5–14 cycle solves. This does not make LAMG+ the fastest solver on this class—these matrices are small and well-separated, so approxChol and CHOLMOD stay ahead—so the value here is robustness, closing LAMG’s documented outlier tail. The mechanism, the per-level two-level factors, and interpolation weight derivation are in App. A.

4.9. Ablation Study. Each choice—the reproduced LAMG choices, the cycle-index calibration, and the two refinements of §3—is load-bearing only on the class that stresses it. Table 4.7 lists their individual effects. The dominant setup cost is the test-vector smoothing, fixed at the cheap $K=4$ everywhere (§3.3); the anisotropy fix comes from caliber-2, not from more test vectors (Table 3.2). The veto and caliber-2 are *complementary* on their target class: on the 256^2 , $\epsilon = 10^{-4}$ grid the veto alone reaches its caliber-1 ceiling ($\text{ACF} \approx 0.46$), caliber-2 alone still stalls at 0.99 (the affinity inversion survives), and only both together give 0.11.

5. Conclusion. LAMG+ is a lean, parameter-free, empirically linear-time graph-Laplacian solver. It reproduces LAMG’s choices, plus two targeted, *local* refinements (§3)—a




FIG. A.1. *Why caliber-2 is a one-dimensional fix. (A) Affinity merges along the strong direction (2:1 x -pairs)—a semicoarsening. (B) The error pointwise GS leaves is smooth in x , oscillatory in y . (C) The semicoarsened grid keeps every y -layer; caliber-1 interpolates the smooth- x mode across each x -pair by a constant ($q \approx 2$, $\rho \approx 0.5$), caliber-2 by a line ($q \rightarrow 1$, $\rho \approx 0.12$).*




FIG. A.2. *The caliber-2 fix (§3.2, on by default), two-grid factor on an anisotropic grid. (A) On locally one-dimensional pockets the second parent drops ρ from the caliber-1 ceiling ≈ 0.49 to ≈ 0.12 , ϵ -robust. (B) The gate is self-targeting: $\approx 96\%$ of fine nodes upgrade where the neighborhood is 1-D and $\approx 0\%$ on the isotropic grid.*

strength-of-connection veto and a selective caliber-2 interpolation. LAMG+ works for structured, web, social, citation, collaboration, road, and grid-aligned anisotropic graphs with no per-graph tuning, and on every family of GKS’s approximate-Cholesky benchmark [13] without fail. LAMG+ and approximate-Cholesky are peers with complementary strengths—approximate-Cholesky faster on social and citation graphs, LAMG+ the fastest and most memory-frugal solver on finite-element/structural matrices. Beyond the linear solver, the multilevel hierarchy could be used for related tasks such as eigenproblems [1, §6.4] and nonlinear systems via the Full Approximation Scheme [28, §8.1].

Toward a parallel LAMG+. The results here are single-threaded; parallelization is the main direction left open. Both the setup and the cycle reduce to scalable sparse primitives—sparse matrix–vector products, maximal independent set, segmented reductions—already mapped to GPUs [31] and scaled to 10^5 cores [32, 33]. The one inherently sequential kernel is the Gauss–Seidel smoother; the standard substitutes are polynomial smoothers [34] or a block-Jacobi Gauss–Seidel that freezes the cross-block coupling and, in our measurements, retains most of serial GS’s smoothing power once the blocks are wide. The open obstacle is distributed coarsening: a low-diameter Laplacian must be repartitioned per coarse level to bound inter-processor communication [36], and whether LAMG+’s purely local affinity and caliber-2 tests preserve its lean operator complexity under a communication-minimizing partition is unresolved. A no-fill aggregation hierarchy is, if anything, friendlier to distribution than a sparsified Cholesky factor whose elimination order is global.

Appendix A. Grid-aligned anisotropy: caliber-2 fix.

On a five-point grid with strong (horizontal) weight 1 and weak (vertical) weight ϵ , the headline inverts the naive expectation: an isotropic grid is easy (ACF = 0.01), *heterogeneous* weights (each edge an independent random factor over three decades) are also easy (ACF = 0.006), but *grid-aligned* anisotropy is the hard case (ACF ≈ 0.34 – 0.39 at $\epsilon \leq 10^{-2}$). The mechanism and the cure are both in the LAMG design. Affinity *does* semicoarsen—it merges along the strong direction, the right *kind* of coarsening given a pointwise GS smoother. The caliber-1 piecewise-constant interpolation recovers half the magnitude of a smooth-in-strong mode, inflating the coarse-operator energy by $q \approx 2$ and pinning the *two-level* factor at $\rho \approx 1 - 1/q \approx 0.5$ [1, §3.4.3]. In a multilevel setting, this slowness compounds across levels: a flat V-cycle ($\gamma = 1$, no recombination) sits at ACF ≈ 0.77 . Increasing γ to 1.5 and employing iterate recombination at every level [1, §3.4.5], we can bound the ACF at 0.39 (Table A.1). We can either accept this limited convergence factor (LAMG) or increase the interpolation caliber to improve it (LAMG+).

A.1. A two-grid local Fourier analysis. The ceiling $\rho \approx \frac{1}{2}$ and its removal by caliber-2 follow from a standard two-grid local Fourier analysis [28, §4.1]. Consider the constant-coefficient infinite grid (strong weight 1 along x , weak weight ϵ along y on $\mathbb{Z}^2$); the Laplacian acts on the error mode $\varphi_\theta(\mathbf{m}) = e^{i\theta \cdot \mathbf{m}}$, $\theta = (\theta_x, \theta_y) \in (-\pi, \pi]^2$. Abbreviate $c = \cos^2 \frac{\theta_x}{2}$ and $s = \sin^2 \frac{\theta_x}{2}$ ($c+s = 1$); then

$$(A.1) \quad \hat{L}(\theta) = 4s + 4\epsilon \sin^2 \frac{\theta_y}{2}.$$

Since $\epsilon \ll 1$, $\hat{L}(\theta)$ is small exactly when $\theta_x \approx 0$ for any θ_y (the algebraically smooth error); affinity therefore semicoarsens 2:1 in x and keeps every y -line (Fig. A.1).

Two-level cycle symbol. Coarsening in x at fixed θ_y , the 2:1 step folds each low mode θ ($\theta_x \in (-\frac{\pi}{2}, \frac{\pi}{2}]$) with its x -harmonic $\theta^* = (\theta_x + \pi, \theta_y)$ (same θ_y); the pair spans a subspace invariant under the two-grid operator, on which the operator is the 2×2 matrix

$$(A.2) \quad \hat{T}(\theta) = \hat{S}(\theta)^\nu \left[I - \hat{P} (\hat{P}^* \hat{L} \hat{P})^{-1} \hat{P}^* \hat{L} \right], \quad \hat{L} = \text{diag}(\hat{L}(\theta), \hat{L}(\theta^*)) =: \text{diag}(\ell_1, \ell_2),$$

with the 2×2 diagonal smoother symbol $\hat{S}(\theta) = \text{diag}(\tilde{S}(\theta), \tilde{S}(\theta^*))$, $\nu = \nu_{\text{pre}} + \nu_{\text{post}} = 3$ Gauss–Seidel sweeps, $\hat{P} \in \mathbb{C}^{2 \times 1}$ the interpolation symbol, and the bracketed coarse-grid correction the $\tilde{L}$ -orthogonal projector off range $\hat{P}$.

Smoother. $\tilde{S}$ is the scalar forward Gauss–Seidel symbol $\tilde{S}(\theta) = 1 - \hat{L}(\theta)/\hat{M}(\theta)$, where $\hat{M}(\theta) = (2+2\epsilon) - e^{-i\theta_x} - \epsilon e^{-i\theta_y}$ is the lower-triangular (already-updated-neighbor) part of L .

Interpolation. Interpolating the fine point midway between two coarse parents has the symbol (up to a unimodular phase that cancels in (A.2))

$$(A.3) \quad \hat{P}_{c1} \propto \begin{pmatrix} \cos \frac{\theta_x}{2} \\ i \sin \frac{\theta_x}{2} \end{pmatrix} \quad (\text{caliber-1, piecewise constant}), \quad \hat{P}_{c2} = \begin{pmatrix} c \\ s \end{pmatrix} \quad (\text{caliber-2, linear}).$$

The decisive difference is the *harmonic* (second) component as $\theta_x \rightarrow 0$: it is $O(\theta_x)$ for caliber-1 but $O(\theta_x^2)$ for caliber-2. Caliber-2 injects asymptotically no energy into the unresolved harmonic of a smooth mode—the order-2 interpolation requirement of [28, §4.3].

Energy factor. The coarse-grid correction in (A.2) is a rank-one projector, so $\det \hat{T} = 0$ and the spectral radius equals the trace:

$$(A.4) \quad \rho(\hat{T}(\theta)) = |\tilde{S}(\theta)^\nu (1 - \beta) + \tilde{S}(\theta^*)^\nu \beta|, \quad \beta = \frac{|\hat{P}_1|^2 \ell_1}{|\hat{P}_1|^2 \ell_1 + |\hat{P}_2|^2 \ell_2} \in [0, 1],$$

where β is the fraction of the interpolant's $\hat{L}$ -energy carried by the smooth harmonic and $q := 1/\beta$ is the local-mode form of the per-node energy ratio (2.3). Substituting (A.3) and letting $\epsilon \rightarrow 0$,

$$(A.5) \quad \beta_{c1} = \frac{1}{2} \quad (q = 2) \quad \text{for every } \theta_x, \quad \beta_{c2} = c \rightarrow 1 \quad (q \rightarrow 1) \quad \text{as } \theta_x \rightarrow 0 :$$

caliber-1 commits exactly half of every interpolant's energy to the wrong harmonic, caliber-2 a vanishing fraction on the smooth modes.

The caliber-1 ceiling. The supremum of (A.4) is attained as $\theta_x \rightarrow 0$, on modes *smooth in* x , where any local smoother is powerless ($|\tilde{S}(\theta)| \rightarrow 1$) while the x -harmonic is fully smoothed ($\tilde{S}(\theta^*) \rightarrow 0$). There $\rho \rightarrow 1 - \beta$, so $\beta_{c1} = \frac{1}{2}$ pins

$$(A.6) \quad \rho_{c1} \xrightarrow{\theta_x \rightarrow 0} 1 - \beta_{c1} = 1 - \frac{1}{q} = \frac{1}{2} \quad (q = 2).$$

This ACF does not improve with more relaxation sweeps per cycle.

Caliber-2 removes it. On those same smooth modes $\beta_{c2} \rightarrow 1$ ($q \rightarrow 1$), so $\rho \rightarrow 1 - \beta \rightarrow 0$: the residual moves to the x -harmonic at $\theta_x = \frac{\pi}{2}$, a genuinely high-frequency, smoothable mode. Evaluating (A.4) under the same forward Gauss–Seidel gives $\sup_{\theta} \rho_{c2} \approx 0.08$ – 0.09 for $\epsilon \leq 10^{-1}$ —textbook two-grid efficiency.

Table A.1 compares caliber-1, the caliber-2 LFA prediction, and the measured caliber-2 two-grid factor against anisotropy strength. The LFA predicts the two-grid factor well; the small excess of the measured value is the least-squares weight's deviation from the ideal $\frac{1}{2}$. Recombination adds a marginal improvement, and the full multilevel solver even beats the two-grid LFA rate for $\epsilon \geq 10^{-2}$.

TABLE A.1

Grid-aligned anisotropy on a 64×64 grid: asymptotic convergence factor by method. Under caliber-1, V-cycle = flat $\gamma=1$ (1,2)-cycle and ML = the full solver ($\gamma=1.5$ + recombination). Under caliber-2, LFA is the local Fourier analysis prediction, 2-level the measured two-grid factor without iterate recombination, +recomb the same with recombination, and ML the full LAMG+ multilevel solver with the caliber-2 refinement. The two-grid columns (LFA, 2-level, +recomb) are measured on a periodic grid to match the infinite-grid analysis; caliber-2 is inert at $\epsilon=1$ (isotropic, 0% upgrade).

ϵ	caliber-1		caliber-2			
	V-cycle	ML	LFA	2-level	+recomb	ML
1	0.74	0.010	—	—	—	0.008
10^{-1}	0.76	0.274	0.08	0.08	0.07	0.007
10^{-2}	0.77	0.335	0.09	0.11	0.09	0.051
10^{-4}	0.77	0.391	0.09	0.12	0.09	0.100

A.2. Fitting the caliber-2 weight. A fine node u whose strong edges reach exactly two coarse aggregates a, b is interpolated by $\phi_u = w \phi_a + (1 - w) \phi_b$; fixing the weights to sum to 1 keeps constants exact. The single weight w is fit by least squares over the *same* K test vectors used for affinity (no extra setup)—a two-parent specialization of the least-squares interpolation of Bootstrap AMG [22]: minimize $\sum_k (x_u^{(k)} - [w x_a^{(k)} + (1 - w) x_b^{(k)}])^2$, whose closed form, with $d_k = x_a^{(k)} - x_b^{(k)}$ and $e_k = x_u^{(k)} - x_b^{(k)}$, is

$$(A.7) \quad w^* = \frac{\sum_k d_k e_k}{\sum_k d_k^2},$$

the regression of u 's profile onto the strong chain's endpoints. We guard $w^* \in [0, 1]$ (a convex combination, so the Galerkin operator remains a Laplacian), snapping near-boundary

values to 0/1; an out-of-range fit signals that u is not truly one-dimensional and is left caliber-1. The gate fires only when the two strong-reachable aggregates are distinct and the fit is admissible—self-targeting at no parameter cost.

A.3. The role of caliber-2. The guard Q of (2.3) already makes the solver robust in every dimension ($\rho \lesssim 1 - 1/Q$, never diverges). What it does not do is make q *small* where anisotropy is persistent: grid-aligned anisotropy and 1-D FE chains sit at the ceiling $q \approx 2$ at every level. Caliber-2 attacks that residue, and only where two conditions coincide: the bound is tight (the strong coupling is a line, so caliber-1 across it is exactly the half-recovered mode) and the cure is cheap (two collinear parents). Both hold iff a node is locally one-dimensional. Hence anisotropic 2-D and 3-D *rods* (a single strong direction) become textbook, while 3-D *plates* (strong coupling in a plane) are locally two-dimensional: the gate sees > 2 strong aggregates and correctly declines, leaving this case to the guard.

Appendix B. Linear-time evaluation of the energy guard. A direct evaluation of (2.3) is superlinear in node degree: recomputing, for each candidate (u, s) pair, the nodal energy by summing over u 's incident edges for each of K test vectors costs $\mathcal{O}(K \deg u)$ per pair, and a node with D admissible neighbors is tested $\mathcal{O}(D)$ times, so the work is $\mathcal{O}(K \deg^2)$ —concentrated on the high-degree hubs that dominate the total. The guard admits an $\mathcal{O}(K)$ -per-candidate form. In one $\mathcal{O}(Km)$ pass accumulate, per node u , the three moments $s_0 = \sum_v w_{uv}$, $s_1^{(k)} = \sum_v w_{uv} x_v^{(k)}$, $s_2^{(k)} = \sum_v w_{uv} (x_v^{(k)})^2$. Then both energies in (2.3) are closed forms: the relaxed (optimal) energy is $\min_y E_u = \frac{1}{2} (s_2^{(k)} - (s_1^{(k)})^2 / s_0)$ and the forced energy is $E_u(\cdot; x_s) = \frac{1}{2} (x_s^2 s_0 - 2x_s s_1^{(k)} + s_2^{(k)})$, so each candidate costs $\mathcal{O}(K)$, independent of degree.

REFERENCES

- [1] O. E. Livne and A. Brandt, *Lean Algebraic Multigrid (LAMG): Fast Graph Laplacian Linear Solver*, SIAM J. Sci. Comput. **34**(4), B499–B522, 2012.
- [2] D. A. Spielman, *Algorithms, Graph Theory, and Linear Equations in Laplacian Matrices*, Proc. ICM, 2010.
- [3] E. G. Boman, B. Hendrickson, and S. Vavasis, *Solving elliptic finite element systems in near-linear time with support preconditioners*, SIAM J. Numer. Anal. **46**(6), 2008.
- [4] S. I. Daitch and D. A. Spielman, *Faster approximate lossy generalized flow via interior point algorithms*, STOC, 2008.
- [5] A. Y. Ng, M. I. Jordan, and Y. Weiss, *On spectral clustering: analysis and an algorithm*, NeurIPS, 2002.
- [6] M. Fiedler, *Algebraic connectivity of graphs*, Czechoslovak Math. J. **23**, 1973.
- [7] T. A. Davis, *Direct Methods for Sparse Linear Systems*, SIAM, 2006 (CHOLMOD).
- [8] A. George, *Nested dissection of a regular finite element mesh*, SIAM J. Numer. Anal. **10**(2), 1973.
- [9] R. J. Lipton, D. J. Rose, and R. E. Tarjan, *Generalized nested dissection*, SIAM J. Numer. Anal. **16**(2), 1979.
- [10] D. A. Spielman and S.-H. Teng, *Nearly-linear time algorithms for graph partitioning, graph sparsification, and solving linear systems*, SIAM J. Comput. **40**(4), 2011.
- [11] J. A. Kelner, L. Orecchia, A. Sidford, and Z. A. Zhu, *A simple, combinatorial algorithm for solving SDD systems in nearly-linear time*, STOC, 2013.
- [12] D. A. Spielman et al., *Laplacians.jl* (approxChol); R. Kyng and S. Sachdeva, *Approximate Gaussian elimination for Laplacians*, FOCS, 2016.
- [13] Y. Gao, R. Kyng, and D. A. Spielman, *Robust and practical solution of Laplacian equations by approximate elimination*, arXiv:2303.00709, 2023 (the approxchol_lap2 solver in

Laplacians.jl).

- [14] C. Chen, T. Liang, and G. Biros, *RCHOL: randomized Cholesky factorization for solving SDD linear systems*, SIAM J. Sci. Comput. **43**(6), 2021.
- [15] A. Brandt, S. F. McCormick, and J. Ruge, *Algebraic multigrid (AMG) for sparse matrix equations*, in *Sparsity and its Applications*, Cambridge, 1984.
- [16] J. W. Ruge and K. Stüben, *Algebraic multigrid*, in *Multigrid Methods*, SIAM Frontiers in Appl. Math. **3**, 1987.
- [17] P. Vaněk, J. Mandel, and M. Brezina, *Algebraic multigrid by smoothed aggregation for second and fourth order elliptic problems*, Computing **56**, 1996.
- [18] Y. Notay, *An aggregation-based algebraic multigrid method*, ETNA **37**, 2010.
- [19] A. Napov and Y. Notay, *An efficient multigrid method for graph Laplacian systems*, ETNA **45**, 2016; and ... II: *Robust aggregation*, SIAM J. Sci. Comput. **39**(5), 2017.
- [20] J. Brannick, Y. Chen, J. Kraus, and L. Zikatanov, *Algebraic multilevel preconditioners for the graph Laplacian based on matching in graphs*, SIAM J. Numer. Anal. **51**(3), 2013.
- [21] B. Lee, *Bringing physics into the coarse-grid selection: approximate diffusion-distance/effective-resistance measures for network analysis and algebraic multigrid for graph Laplacians*, Numer. Linear Algebra Appl. **31**(2), 2024.
- [22] A. Brandt, J. Brannick, K. Kahl, and I. Livshits, *Bootstrap AMG*, SIAM J. Sci. Comput. **33**(2), 2011.
- [23] D. Ron, I. Safro, and A. Brandt, *Relaxation-based coarsening and multiscale graph organization*, Multiscale Model. Simul. **9**(1), 2011.
- [24] I. Luz, M. Galun, H. Maron, R. Basri, and I. Yavneh, *Learning algebraic multigrid using graph neural networks*, ICML, 2020.
- [25] I. Koutis, G. L. Miller, and D. Tolliver, *Combinatorial preconditioners and multilevel solvers for problems in computer vision and image processing*, Comput. Vis. Image Underst. **115**(12), 2011.
- [26] N. Bell, L. N. Olson, J. Schroder, and B. Southworth, *PyAMG: algebraic multigrid solvers in Python*, J. Open Source Softw. **8**(87), 2023.
- [27] V. E. Henson and U. M. Yang, *BoomerAMG: a parallel algebraic multigrid solver and preconditioner*, Appl. Numer. Math. **41**(1), 2002.
- [28] A. Brandt, *Multigrid Techniques: 1984 Guide with Applications to Fluid Dynamics*, GMD-Studie 85, 1984; reissued SIAM, 2011.
- [29] O. E. Livne, *LAMG+: open-source Julia implementation with all benchmark scripts*, <https://github.com/orenlivne/lamgplus>, 2025.
- [30] M. A. Christie and M. J. Blunt, *Tenth SPE comparative solution project: a comparison of upscaling techniques*, SPE Reservoir Eval. & Eng. **4**(4), 308–317, 2001 (the SPE10 benchmark; public permeability field).
- [31] N. Bell, S. Dalton, and L. N. Olson, *Exposing fine-grained parallelism in algebraic multigrid methods*, SIAM J. Sci. Comput. **34**(4), C123–C152, 2012.
- [32] A. H. Baker, R. D. Falgout, T. V. Kolev, and U. M. Yang, *Scaling hypre's multigrid solvers to 100,000 cores*, in *High-Performance Scientific Computing*, Springer, 2012.
- [33] M. Naumov et al., *AmgX: a library for GPU accelerated algebraic multigrid and preconditioned iterative methods*, SIAM J. Sci. Comput. **37**(5), S602–S626, 2015.
- [34] M. Adams, M. Brezina, J. Hu, and R. Tuminaro, *Parallel multigrid smoothing: polynomial versus Gauss–Seidel*, J. Comput. Phys. **188**(2), 593–610, 2003.
- [35] E. Cuthill and J. McKee, *Reducing the bandwidth of sparse symmetric matrices*, in Proc. 24th National Conference of the ACM, 1969, pp. 157–172.
- [36] G. Karypis and V. Kumar, *A fast and high quality multilevel scheme for partitioning irregular graphs*, SIAM J. Sci. Comput. **20**(1), 359–392, 1998.
- [37] G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, and T.-Y. Liu, *LightGBM: a highly efficient gradient boosting decision tree*, in *Advances in Neural Information Processing*

Systems (NeurIPS) **30**, 2017.